

Homework 7

1. A *dominating set* in an undirected graph $G = (V, E)$ is a subset of nodes $U \subset V$ such that every node is either in U or adjacent to (that is, has an edge to) U . The problem of finding the *smallest* dominating set is of practical importance, but is NP-complete; no efficient algorithm is known for it. We'll use randomization to construct a pretty good solution, in the case where each node of G has degree d .

$U = \emptyset$

For each $v \in V$:

 Add v to U with probability p

- (a) Pick any node $v \in V$. We will say that v is *covered* if either $v \in U$ or if v is adjacent to U . What is the probability that v is not covered, as a function of p and d ?
 - (b) Show that for $p = (\ln 2n)/(d+1)$, where $n = |V|$, the expected number of uncovered nodes is less than $1/2$. *Hint:* use the inequality $1 - x \leq e^{-x}$.
 - (c) For this setting of p , what is the expected size of U ?
 - (d) **(Challenging)** Show that for this setting of p , the probability that U is a valid dominating set is at least $1/2$. That is, the algorithm fails with probability less than $1/2$.
 - (e) Assuming part (d) is true, how can the probability of failure be reduced to δ ?
2. Barbara Smith is interviewing candidates to be her secretary, one at a time. After each interview she is able to determine the true competence level of the candidate (which can be thought of as some positive real number). However, she needs to make a spot decision whether or not to hire a candidate, before interviewing the remaining ones. If the candidates appear in random order, what is a good hiring strategy for Barbara?
- Suppose there are n candidates. Barbara decides to interview the first r candidates, noting down their scores but not hiring any of them. Let s be the largest score she records during this time. Then, she starts interviewing the remaining $n - r$ candidates, and hires the first one who scores more than s (if none of them do, then she simply picks the last candidate).
- (a) Show that the probability that Barbara ends up with the best secretary is at least $r(n - r)/n^2$.
Hint: Consider the event that the second-best secretary is in the first group of r people, while the best secretary is the second group.
 - (b) What is a good setting for r , and what is the probability of success in this case?
3. Suppose you have access to two different algorithms for testing whether a number is prime. The first is a randomized algorithm \mathcal{A} that runs in time $T(n)$ and has the following behavior:

- If x is prime, $\mathcal{A}(x)$ always says “prime”
- If x is not prime, $\mathcal{A}(x)$ says “not prime” with probability $1/2$; otherwise it says “prime”

Algorithm \mathcal{B} is deterministic, runs in time $100T(n)$, and always returns the correct answer.

You wish to create an algorithm that is always correct and is as efficient as possible.

- (a) Specify your algorithm.
 - (b) What is the expected running time of your algorithm when the input is prime?
 - (c) When the input is not prime?
4. Suppose you have a linked list of n elements *in sorted order*. Here's the obvious way to look up an element x , given a pointer to the first node in the list:

```

Function lookup( $L, x$ )
  if  $L$  is NULL: return NULL
  if value( $L$ ) =  $x$ : return  $L$ 
  if value( $L$ ) >  $x$ : return NULL
  return lookup(next( $L$ ),  $x$ )

```

Notation: given the pointer P to a node, $\text{value}(P)$ and $\text{next}(P)$ are the value in the node and the pointer to the next node.

- (a) Suppose an element of L is chosen at random. What is (roughly) the expected time (in terms of n) to look up this element?
 - (b) One way to get faster lookups is to add a second pointer $\text{jump}(P)$ that points to the node k steps down the list from P (that is, following a “jump” pointer is like following k “next” pointers). The procedure above is then altered by adding the following line just before the final return statement:


```

      if value(jump( $L$ )) ≤  $x$ : return lookup(jump( $L$ ),  $x$ )
      
```

 With this addition, the data structure is called a *skip list*. Now suppose an element of L is chosen at random. What is, roughly, the expected time to look it up, as a function of n and k ?
 - (c) What is the best choice of k ?
5. For cryptographic applications, it is necessary to generate random n -bit prime numbers. One way to do this is to repeatedly generate random n -bit numbers until you find one that is prime. Suppose that we check primality using an algorithm \mathcal{A} with the following behavior on n -bit numbers:

- If x is prime, $\mathcal{A}(x)$ always says “prime”
- If x is not prime, $\mathcal{A}(x)$ says “prime” with probability at most $1/(10^6 n)$; otherwise “not prime”

The overall scheme for generating a random n -bit prime then looks like this:

- Repeat:
 - Generate a random n -bit number x .
 - If $\mathcal{A}(x) = \text{“prime”}$: return x and halt.

A useful fact from number theory: a random n -bit number has approximately a c/n probability of being prime, for some small constant c . To keep things simple, assume this probability is exactly $1/n$.

- (a) Give an upper bound on the probability that the algorithm ends up returning a number that is not prime. *Hint*: use Bayes’ rule.
 - (b) What is the expected running time of the algorithm, assuming \mathcal{A} runs in time $O(n^3)$?
6. You are dealt ten cards, one at a time, from the top of a randomly shuffled deck.
- (a) What is the probability that the tenth card you get is an ace?

- (b) What is the probability that the tenth card is an ace, given that the third card is an ace?
- (c) What is the probability that the third card is an ace, given that the tenth card is an ace?
7. Pick a random permutation of $(1, 2, \dots, n)$. Let X_i be the number that ends up in the i th position. For instance, if the permutation is $(3, 2, 4, 1)$ then $X_1 = 3$, $X_2 = 2$, $X_3 = 4$, and $X_4 = 1$.
- (a) What is the expected number of positions at which $X_i \neq i$?
- (b) What is the expected number of positions at which $X_i = i + 1$?
- (c) What is the expected number of positions at which $X_i \geq i$?
- (d) What is the expected number of positions at which $X_i > \max(X_1, \dots, X_{i-1})$?
8. A set of n people are lined up along a wall in a random order. Among them are Alice, Bob, and Chet.
- (a) What is the probability that Alice appears somewhere to the left of Bob, and Bob appears somewhere to the left of Chet?
- (b) What is the expected number of people between Alice and Bob?
9. In class, we saw how to generate fair coin flips given only a coin of unknown bias. But suppose we have a coin whose bias, p , we *know*. Can we then do better, in terms of using fewer coin flips in expectation? Here's a potential scheme which assumes (without loss of generality) that $p < 1/2$. In fact, this scheme can generate a coin flip of any desired bias q , given the coin of bias p . To get a fair coin flip, call it with $q = 1/2$.

```

Function simulate(q)
  If  $q < 1/2$ :
    Return switch(simulate(1 - q)) // 'switch' converts H to T and T to H
  Flip the coin of bias  $p$ 
  If it turns up heads:
    Return  $H$ 
  else:
    Return simulate((q - p)/(1 - p))

```

- (a) Can you see why this algorithm will output H with probability exactly q ? To show this, assume the recursive calls work correctly, and show that the top-level algorithm will do the right thing. Separately consider the two cases where $q < 1/2$ and $q \geq 1/2$.
- (b) What is the expected number of times the biased coin is flipped?
10. **(Challenging)** *From Las Vegas to Monte Carlo.* Suppose you have an algorithm \mathcal{A} for your problem that always returns the correct answer, but takes different amounts of time each time it runs. Its *expected* time on an input of size n is $T(n)$. What you would prefer, however, is an algorithm that *always* finishes in time $O(T(n))$, but may have up to a 5% probability of returning the wrong answer. Show how to construct such an algorithm from \mathcal{A} . *Hint:* Start by showing that for any $c > 1$, the probability that \mathcal{A} takes longer than $cT(n)$ is at most $1/c$. Then use a timer while running \mathcal{A} !

Answer key 1 (a) $(1 - p)^{d+1} (c) (n \ln 2n)/(d + 1) (e)$ Repeat the algorithm $\log 1/\delta$ times, and return the smallest valid dominating set found (if any). 2. Choosing $r = n/2$ gives at least a $1/4$ probability of success. 3 (a) If $\mathcal{A}(x) =$ "not prime" return "not prime" else return $\mathcal{B}(x)$ (b) $101T(n)$ (c) $51T(n)$ (d) roughly $n/2$ (e) roughly $n/2k + k/2$ (f) \sqrt{n} . 5 (a) 10^{-6} (b) $O(n^{1/4})$. 6 (a) $1/13$ (b) $1/17$ (c) $1/17$. 7 (a) $u - 1$ (b) $1 - u$ (c) $u/(1 + u)$ (d) approximately $\ln n$. 8 (a) $1/6$ (b) $(n - 2)/3$. 9 (b) $1/p$. 10. Run \mathcal{A} for $20T(n)$ time steps; if it is still going, halt and return some generic answer.